

Introduction to Theano

Ian Goodfellow

Outline

- Overview of library (3 min)
- Building expressions (30 min)
- Compiling and running expressions (30 min)
- Modifying expressions (25 min)
- Debugging (30 min)
- Citing theano (2 min)

Overview of Library

- Theano is many things
 - Language
 - Compiler
 - Python library

Overview

- Theano language:
 - Operations on scalar, vector, matrix, tensor, and sparse variables
 - Linear algebra
 - Element-wise nonlinearities
 - Convolution
 - Extensible

Overview

- Using Theano:

- define expression $f(x, y) = x + y$

- compile expression

```
int f(int x, int y) {  
    return x + y;  
}
```

- execute expression

```
>>> f(1,2)  
3  
>>> █
```

Building expressions

- Scalars
- Vectors
- Matrices
- Tensors
- Reductions
- Dimshuffle

Scalar math

```
from theano import tensor as T
x = T.scalar()
y = T.scalar()
z = x + y
w = z * x
a = T.sqrt(w)
b = T.exp(a)
c = a ** b
d = T.log(c)
```

Vector math

```
from theano import tensor as T
x = T.vector()
y = T.vector()
# Scalar math applied elementwise
a = x * y
# Vector dot product
b = T.dot(x, y)
# Broadcasting
c = a + b
```


Matrix math

```
from theano import tensor as T
x = T.matrix()
y = T.matrix()
a = T.vector()
# Matrix-matrix product
b = T.dot(x, y)
# Matrix-vector product
c = T.dot(x, a)
```

Tensors

- Dimensionality defined by length of “broadcastable” argument
- Can add (or do other elemwise op) on two tensors with same dimensionality
- Duplicate tensors along broadcastable axes to make size match

```
from theano import tensor as T
```

```
tensor3 =
```

```
T.TensorType(broadcastable=(False, False,  
False), dtype='float32')
```

```
x = tensor3()
```

Reductions

```
from theano import tensor as T
tensor3 =
T.TensorType(broadcastable=(False,
False, False), dtype='float32')
x = tensor3()
total = x.sum()
marginals = x.sum(axis=(0, 2))
mx = x.max(axis=1)
```

Dimshuffle

```
from theano import tensor as T
tensor3 =
T.TensorType(broadcastable=(False,
False, False), dtype='float32')
x = tensor3()
y = x.dimshuffle((2, 1, 0))
a = T.matrix()
b = a.T
# Same as b
c = a.dimshuffle((0, 1))
# Adding to larger tensor
d = a.dimshuffle((0, 1, 'x'))
e = a + d
```

Exercises

- Work through the “01_building_expressions” directory now

Compiling and running expressions

- `theano.function`
- shared variables and updates
- compilation modes
- compilation for GPU
- optimizations

theano.function

```
>>> from theano import tensor as T
>>> x = T.scalar()
>>> y = T.scalar()
>>> from theano import function
>>> # first arg is list of SYMBOLIC inputs
>>> # second arg is SYMBOLIC output
>>> f = function([x, y], x + y)
>>> # Call it with NUMERICAL values
>>> # Get a NUMERICAL output
>>> f(1., 2.)
array(3.0)
```

Shared variables

- It's hard to do much with purely functional programming
- “shared variables” add just a little bit of imperative programming
- A “shared variable” is a buffer that stores a numerical value for a theano variable
- Can write to as many shared variables as you want, once each, at the end of the function
- Modify outside function with `get_value` and `set_value`

Shared variable example

```
>>> from theano import shared
>>> x = shared(0.)
>>> from theano.compat.python2x import OrderedDict
>>> updates = OrderedDict()
>>> updates[x] = x + 1
>>> f = function([], updates=updates)
>>> f()
[]
>>> x.get_value()
1.0
>>> x.set_value(100.)
>>> f()
[]
>>> x.get_value()
101.0
```

Which dict?

- Use `theano.compat.python2x.OrderedDict`
- Not `collections.OrderedDict`
 - This isn't available in older versions of python, and will limit the portability of your code
- Not `{}` aka dict
 - The iteration order of this built-in class is not deterministic (thanks, Python!) so if Theano accepted this, the same script could compile different C programs each time you run it

Compilation modes

- Can compile in different modes to get different kinds of programs
- Can specify these modes very precisely with arguments to `theano.function`
- Can use a few quick presets with environment variable flags

Example preset compilation modes

- `FAST_RUN`: default. Spends a lot of time on compilation to get an executable that runs fast.
- `FAST_COMPILE`: Doesn't spend much time compiling. Executable usually uses python instead of compiled C code. Runs slow.
- `DEBUG_MODE`: Adds lots of checks. Raises error messages in situations other modes regard as fine.

Compilation for GPU

- Theano only supports 32 bit on GPU
 - CUDA supports 64 bit, but is slow
 - T.fscalar, T.fvector, T.fmatrix are all 32 bit
 - T.scalar, T.vector, T.matrix resolve to 32 bit or 64 bit depending on theano's floatX flag
 - floatX is float64 by default, set it to float32
- Set device flag to gpu (or a specific gpu, like gpu0)

Optimizations

- Theano changes the symbolic expressions you write before converting them to C code
- It makes them faster
 - $(x+y)+(x+y) \rightarrow 2(x+y)$
- It makes them more stable
 - $\exp(a)/\exp(a).\text{sum}() \rightarrow \text{softmax}(a)$

Optimizations

- Sometimes optimizations discard error checking and produce incorrect output rather than an exception

```
>>> x = T.scalar()  
>>> f = function([x], x/x)  
>>> f(0.)  
array(1.0)
```

Exercises

- Work through the
“02_compiling_and_running” directory
now

Modifying expressions

- The grad method
- Variable nodes
- Types
- Ops
- Apply nodes

The grad method

```
>>> x = T.scalar('x')
>>> y = 2. * x
>>> g = T.grad(y, x)
>>> from theano.printing import min_informative_str
>>> print min_informative_str(g)
```

- A. Elemwise{mul}
 - B. Elemwise{second,no_inplace}
 - C. Elemwise{mul,no_inplace}
 - D. TensorConstant{2.0}
 - E. x
 - F. TensorConstant{1.0}
- <D>

Theano Variables

- A *Variable* is a theano expression
- Can come from T.scalar, T.matrix, etc.
- Can come from doing operations on other Variables
- Every Variable has a type field, identifying its Type
 - e.g. `TensorType((True, False), 'float32')`
- Variables can be thought of as *nodes* in a *graph*

Ops

- An Op is any class that describes a mathematical function of some variables
- Can call the op on some variables to get a new variable or variables
- An Op class can supply other forms of information about the function, such as its derivatives

Apply nodes

- The *Apply* class is a specific instance of an application of an Op
- Notable fields:
 - op: The Op to be applied
 - inputs: The Variables to be used as input
 - outputs: The Variables produced
- Variable.owner identifies the Apply that created the variable
- Variable and Apply instances are *nodes* and owner/inputs/outputs identify *edges* in a Theano graph

Exercises

- Work through the “03_modifying” directory now

Debugging

- `DEBUG_MODE`
- `compute_test_value`
- `min_informative_str`
- `DebugPrint`
- `Accessing the FunctionGraph`

compute_test_value

```
>>> from theano import config
>>> config.compute_test_value = 'raise'
>>> x = T.vector()
>>> import numpy as np
>>> x.tag.test_value = np.ones((2,))
>>> y = T.vector()
>>> y.tag.test_value = np.ones((3,))
>>> x + y
...
ValueError: Input dimension mis-match.
(input[0].shape[0] = 2, input[1].shape[0] = 3)
```


min_informative_str

```
>>> x = T.scalar()
>>> y = T.scalar()
>>> z = x + y
>>> z.name = 'z'
>>> a = 2. * z
>>> from theano.printing import min_informative_str
>>> print min_informative_str(a)
```

- A. Elemwise{mul,no_inplace}
- B. TensorConstant{2.0}
- C. z

debugprint

```
>>> from theano.printing import debugprint
>>> debugprint(a)
Elemwise{mul,no_inplace} [@A] "
|TensorConstant{2.0} [@B]
|Elemwise{add,no_inplace} [@C] 'z'
|<TensorType(float64, scalar)> [@D]
|<TensorType(float64, scalar)> [@E]
```

Accessing a function's fgraph

```
>>> x = T.scalar()
>>> y = x / x
>>> f = function([x], y)
>>> debugprint(f.maker.fgraph.outputs[0])
DeepCopyOp [@A] "
|TensorConstant{1.0} [@B]
```

Exercises

- Work through the “04_debugging” directory now

Citing Theano

- Please cite both of the following papers in all work that uses Theano:
 - Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Bergstra, James, Goodfellow, Ian, Bergeron, Arnaud, Bouchard, Nicolas, and Bengio, Yoshua. Theano: new features and speed improvements. Deep Learning and Unsupervised Feature Learning NIPS 2012 Workshop, 2012.
 - Bergstra, James, Breuleux, Olivier, Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Desjardins, Guillaume, Turian, Joseph, Warde-Farley, David, and Bengio, Yoshua. Theano: a CPU and GPU math expression compiler. *In Proceedings of the Python for Scientific Computing Conference (SciPy)*, June 2010. Oral Presentation.

Example acknowledgments

We would like to thank the developers of Theano~
\citep{bergstra+al:2010-scipy,Bastien-Theano-2012},
Pylearn2~\citep{pylearn2_arxiv_2013}. We would also like
to thank NSERC, Compute Canada, and Calcul Qu'\ebec
for providing computational resources.